

# HOMEWORK #1:

## *Linear Logistic Classification*


**Due Date:** Thursday, December 9th, 11:59:59pm

### **Problem:**

Write a single-neuron **Logistic Classifier**, a program that will learn to classify inputs from a collection of examples.

### **Input:**

The file `train.txt` contains the training data for this classification problem. The first two columns are the **inputs** of each example ( from 0 to 100 ). The third column is the classification of the example inputs. ( 0 or 1 ). The file `valid.txt` contains the **validation** data for this classification problem. You can visualize the training data in the following plot:



## Gradient Descent:

You will implement your Logistic Classifier using **incremental gradient descent**, and will use the Sum-of-Squares as the error measurement for your classifier. Your objective is then to minimize the function:

$$SSE(E) = \sum_{e \in E} (y(e) - \hat{y}(e))^2$$

where  $E$  is the set of examples,  $y(e)$  is the class value for an example  $e$ , and  $\hat{y}(e)$  is the output of the classifier, given by

$$z = \sum_{i=0}^n (w_i * x_i(e))$$
$$\hat{y}(e) = \text{sig}(z)$$

Where  $\text{sig}()$  is the sigmoid function given by

$$\text{sig}(z) = 1 / (1 + e^{-z})$$

Recall also the derivative of the sigmoid function:

$$\text{sig}'(z) = \text{sig}(z) * (1 - \text{sig}(z))$$

If the error for a single example  $e$  is:

$$(y(e) - \hat{y}(e))^2$$

Then the *partial derivative* of this error with respect to a single weight  $w_i$  is:

$$\begin{aligned} & (\partial \text{Error} / \partial \hat{y}(e)) * (\partial \hat{y}(e) / \partial z) * (\partial z / \partial w_i) \\ & = \\ & -2 * (y(e) - \hat{y}(e)) * \hat{y}(e) * (1 - \hat{y}(e)) * x_i(e) \end{aligned}$$

The constant 2 is latter absorbed into the learning rate  $\eta$  (eta).

Notice that we are applying the sigmoid function as a **squash function**. This is because we are doing *classification*. Also notice that this problem has **two** inputs, which means that the classifier is going to be learning **three** weights.

You shall implement **batch** gradient descent. This means that you will be updating the learner's weights after each batch. Initialize the weights randomly with real numbers **in the range [-2..2]**. The **learning rate  $\eta$**  (eta), the **number of iterations** and the **batch size** is **left for you to decide**. Experiment with different values and search for one that yields good results.

## Validation:

After performing gradient descent on the training data, you should evaluate the performance of your learner against the validation data by computing the sum-of-squares error between the validation data and your learner's predictions. Do **not** use the validation data to train!.

## Submission Guidelines:

You will submit through **Canvas**:

Your submission should consist of the following components:

1. Your **program** files.- Submit all necessary files. Your main program file should be called 'learner1.X' where X is the extension of whatever programming language you are using.
2. A **text** file, called 'learner1output.txt', in which you report on the training run that achieved the best results. This file should **strictly** follow the format shown below.

### Report file format:

The first line of the report text file is the values of the learned weights, space separated. The second line is the **Sum of Squares** error of the learned weights against the validation data. It is important that you follow the format because **these lines will be read by the auto-grader** to evaluate your program.

Sample 'learner1output.txt'

```
3.1415 42.0 1.0
1237491.2831
```

```
CS-5001: HW#1
Programmer: Philip J. Fry
```

```
TRAINING:
Using learning rate eta = 0.001
Using 25000 iterations.
```

```
LEARNED:
w0 = 3.1415
w1 = 42.0
w2 = 1.0
```

VALIDATION

Sum-of-Squares Error = 1237491.2831

## Pseudocode:

PROCEDURE Logistic\_Regression\_Learner

GIVEN:

Ex[0..n] : examples, each a  $\langle X_0 \dots X_i, Y \rangle$  tuple.  
( remember, in every example,  $X_0 = 1$  )

LOCAL:

w[0..i] : weights.  
Randomize w[0..i].

REPEAT no\_iterations

FOR EACH example e =  $\langle X_0 \dots X_i, Y \rangle$  DO  
    compute yCap ( from w[0..i] and  $X_0 \dots X_i$  )  
    delta :=  $( Y - YCap ) * YCap * ( 1 - Ycap )$   
    FOR EACH weight w[k] DO  
        update w[k] by  $Eta * delta * X_k$

OUTPUT w[0..i]

**END.**